

Programming Assignment #2

Author: Yu Sheng Lin

Instructor: Wei Chao Chen

April 26, 2018

1 Goals

You have to

1. Learn how to use the Thrust library.
2. Learn how to implement some parallel algorithms on GPU.

in this assignment.

2 Requirements

This is a two-part fixed assignment. You are asked to implement the same functionality in both part. However, in part I you are only allowed to use existing libraries while in part II the efficiency of your code is also considered during grading.

In this assignment you are required to count the position of each character inside the word it belongs to. Listing 1. provides a sample input and output.

```
1 gpu qq a hello sonoda (input)
2 123012001012345000123456 (output)
```

Listing 1: Example: Count the Position in Words.

You may easily come up with an $O(n)$ sequential algorithm and an $O(nk)$ parallel algorithm, where n is the length of the input and k is the maximum length of a word.

The input is generated in a pseudo-random manner, while we will use a different random seed during grading. You can assume that $k = 500$, $n \approx 4 \times 10^7$ and the input only contains characters [a-z] and we use linebreak '\n' as the spaces.

You have to implement a function whose signature is Listing 2.. All pointers are device pointers and `text_size` is the n . It should also be noticed that `gridDim.x` cannot exceed $2^{17} = 131072$ if you don't use `-arch sm_30` (or higher) ¹ compile flag.

2.1 Part I: Using the Thrust Library (40pts)

Thrust is a useful API including many common parallel computing patterns, and in this part you should only include

¹We assume that you have a GPU newer than Kepler architecture.

```
1 void CountPosition1(const char *text, int *pos, int text_size);
2 void CountPosition2(const char *text, int *pos, int text_size);
```

Listing 2: The function signature of part I.

- Declaration of native and `thrust::*` types,
- A few `structs` with call operator and
- `thrust::*` and native CUDA functions like `cudaFree` (namely `__global__` functions are not allowed)

in `CountPosition1`.

Here are some hints:

- You will need the document <https://thrust.github.io/doc/modules.html>
- I have already included some necessary headers.
- If you write more than 10 lines, then you are probably wrong.

2.2 Part II: Implement Your Own Kernel (40pts+15pts bonus)

In part II, you have to implement the same functionality from scratch, and **no external API and library is allowed**. We provide some hints about the $O(n \ln k)$ algorithm in a separate PDF while it's not the only solution, and you can decide whether to read them by yourself.

According to our experience, when $k = 500$, $O(n \ln k)$ and $O(nk)$ might have comparable speed. To achieve the best efficiency and outperform your classmate, we also challenge you to implement both and compare them if possible. According to the execution time of your program, at most 15pts bonus will be given.

3 Submission

- The submission deadline is 2017/5/10 23:59 (Fri.).
- The efficiency of part I will NOT be considered during grading, but if you break the rules listed in part 2.1, you will get 0pt.
- The efficiency part II will also be considered during grading, and pure CPU implementations would be disqualified.
- You can only modify `lab2/counting.cu`, and we will only copy this file from your repo.
- The compile flags are `--std=c++11 -O2 -arch sm_50`, and the environment is ArchLinux with CUDA 9.1. Although we will test your code on a Linux machine with one GTX 970, you should not use platform dependent libraries.
- Do not hack the answer by inspecting the memory or modifying the stack. This is a course about GPGPU programming, not assembly programming.
- Please also refer to assignment #0 for more details.